

Hierarchical Semantic Memory Architecture

A proposed memory and hardware co-design for long-horizon agentic AI

Core thesis

Long-horizon agents need memory that scales with accumulated understanding, not only with token history.

HSMA proposes a multi-level hierarchy that can demote raw KV into compressed neural state, semantic summaries, concept nodes, and a knowledge graph - while preserving pointers for selective reconstruction.

The opportunity is a new systems layer between inference, retrieval, and storage: semantic memory management for agentic AI.

Prepared for: Senior Technical Leadership


Prepared by: Oluwole Jaiyeoba

Document type: Internal concept proposal / architecture whitepaper

Version: 2.0

HSMA: Memory as a Hierarchy of Understanding

Progressive demotion preserves pointers while reducing storage cost and preserving future reconstructability.



Core operation: Tokens -> Facts -> Concepts -> Understanding

Critical property: each abstract level maintains pointers to the lower-fidelity or full-fidelity source state.

Figure 1. HSMA reframes LLM memory as progressive semantic abstraction with downward pointers.

1. Executive One-Page Summary

Large language models are moving from short-lived inference requests toward persistent agents that plan, use tools, collaborate, and accumulate state across long time horizons. In these workloads, memory pressure is no longer a secondary issue. It becomes a first-order systems bottleneck that affects cost, latency, accuracy, and deployability.

Proposal in one sentence

Build a Hierarchical Semantic Memory Architecture (HSMA) that continuously converts raw token/KV history into compressed KV, summaries, concepts, and knowledge graphs - while keeping pointers that allow the system to reconstruct only the context it needs.

Dimension	Executive summary
Problem	Current systems largely retain, compress, summarize, or retrieve history. They do not maintain a unified hierarchy from neural state to semantic understanding.
Architecture	HSMA: Raw KV -> Compressed KV -> Semantic Summaries -> Concept Nodes -> Knowledge Tree / Graph.
Novel contribution	A semantic demotion and selective reconstruction layer that links KV-cache management, RAG, summaries, and graph memory into one pointer-preserving hierarchy.
Hardware opportunity	Semantic Memory Controller, Knowledge Cache, Graph Accelerator, KV Demotion Engine, Reconstruction Engine, and longer-term Knowledge Processing Unit.

The strategic hypothesis is that future agentic systems will need a memory abstraction layer comparable in importance to virtual memory, cache hierarchy, and storage hierarchy in classical systems. The difference is that the movement policy is based not only on recency or address locality, but also on semantic utility, future relevance, and reconstructability.



2. Problem Statement: Agentic AI Changes the Memory Problem

Traditional inference workloads are bounded: a prompt arrives, tokens are generated, and the request ends. Agentic systems are different. They can run across sessions, use tools, write code, coordinate sub-agents, remember user preferences, and revisit decisions made thousands or millions of tokens earlier.

CURRENT SCALING PATTERN

More interaction -> more tokens -> more KV state -> more memory movement -> higher cost and latency.

The KV cache is especially important because it avoids recomputing prior key/value states during autoregressive generation, but its size grows with context length and batch size. Recent KV-cache research explicitly identifies the cache as a major inference memory bottleneck and proposes eviction, sparsity, and quantization techniques to reduce footprint [6, 8, 9, 10].

Reframing

The core question should shift from "How do we store more tokens?" to "How do we preserve useful understanding at the lowest necessary fidelity?"

3. Prior Art Landscape and Limitations

HSMA does not replace the existing research landscape. It attempts to connect several strong but currently separate lines of work into a single systems abstraction.

Approach	Primary unit	Representative work	Strength	Gap relative to HSMA
KV compression / eviction	Neural state	H2O, KIVI, SnapKV, PyramidKV	Reduces KV footprint and bandwidth	Does not create a semantic hierarchy or future reconstruction model
Long-context memory models	Hidden state / memory	Transformer-XL, Compressive Transformer	Extends temporal context	Mostly model-internal, not a full agent memory OS
RAG / RETRO	Text chunks	RAG, RETRO	Scalable retrieval from external memory	Usually flat retrieval; weak persistent relational structure
Agent memory	Notes / summaries	MemGPT, generative agents	Practical long-term memory management	Mostly text-level memory; limited KV integration
GraphRAG	Entities / graph	Microsoft GraphRAG	Structured retrieval and community summaries	Typically disconnected from neural cache and reconstruction path


The key observation is not that any one prior-art line is insufficient. The gap is that the system lacks a unified transition path from raw neural state to compressed state to semantic structures and back again when needed.

4. Proposed Architecture: Hierarchical Semantic Memory Architecture

HSMA introduces a multi-resolution memory hierarchy for long-horizon agentic systems. Each level stores the same underlying information at a different level of fidelity, cost, and reconstructability.

HSMA: Memory as a Hierarchy of Understanding

Progressive demotion preserves pointers while reducing storage cost and preserving future reconstructability.



Core operation: Tokens -> Facts -> Concepts -> Understanding

Critical property: each abstract level maintains pointers to the lower-fidelity or full-fidelity source state.

Figure 2. The HSMA hierarchy from raw neural state to structured understanding.

Level	Representation	Primary role	Storage / compute profile
0	Raw KV cache	Maximum fidelity for active reasoning	High HBM footprint; fastest attention path
1	Compressed KV	Retain neural state at lower cost	Quantized, pruned, clustered, or offloaded
2	Semantic summaries	Compact textual or latent meaning	Cheap to store; may lose detail
3	Concept nodes	Entities, events, decisions, goals, constraints	Structured metadata and references
4	Knowledge tree / graph	Relationships and navigable understanding	Lowest cost; highest abstraction

Design principle

Higher levels do not merely summarize lower levels. They maintain typed pointers to source tokens, raw text, compressed KV, raw KV segments, timestamps, task state, and dependency relationships.

5. What Is Existing vs. What Is Novel

A credible architecture proposal should avoid claiming novelty for established ideas. HSMA builds on strong prior art, but the proposed contribution is the integration and systems co-design.

Category	Examples	Assessment
Existing prior art	KV compression, RAG, agent memory, graph retrieval, OS-inspired context management	Do not claim these as new. They are foundation technologies.
Incremental extensions	Better memory policies, graph-enhanced retrieval, compressed KV offload, summary hierarchies	Valuable but likely incremental unless integrated into a stronger architecture.
Proposed novel contribution	Semantic demotion + pointer-preserving hierarchy + selective reconstruction + hardware/runtime co-design	This is the investment-worthy system abstraction.

The strongest novelty claim is not "we compress KV" or "we build a graph." The stronger claim is: we manage neural, textual, and symbolic memory as a unified hierarchy and introduce runtime/hardware mechanisms for demotion, promotion, and selective reconstruction.

6. Architectural Walkthrough

Consider an agent that has discussed a product architecture, constraints, experiments, and unresolved decisions over many sessions. A flat transcript or vector store can retrieve matching chunks. HSMA instead maintains a navigable structure of the conversation and its reasoning state.

Step	Operation	Example
1. Observe	Process incoming tokens and produce raw KV	Conversation contains: "Paris is the capital of France."
2. Extract	Identify facts, entities, dependencies, and task relevance	Entity: France. Relation: capital_of(Paris, France).
3. Demote	Move cold state from raw KV into compressed KV or summaries	Old geographic fact no longer needs full HBM residency.
4. Link	Attach pointers from concept nodes to summaries, source text, and KV segments	Concept node France points to source span and compressed KV block.
5. Retrieve	When a question arrives, traverse concepts before loading detailed state	Question: "Which country did we discuss?" -> Geography -> Europe -> France.
6. Reconstruct	Rebuild only the relevant context branch	Load summary, source span, or replay tokens to regenerate KV as needed.


Important technical nuance

The architecture does not assume that semantic summaries can perfectly replace KV. It assumes multiple representations can coexist, and that fidelity should be allocated based on predicted future utility and reconstruction cost.

7. Proposed Hardware and Runtime Architecture

The near-term implementation should be software-first: a runtime, metadata store, KV compression layer, graph store, and reconstruction policy. The long-term hardware opportunity is to identify which semantic-memory operations become common enough, expensive enough, and structured enough to justify acceleration.

HSMA Hardware / Runtime Co-Design View



Near-term implementation: software runtime + GPU kernels. Long-term research: hardware-assisted semantic movement and

Figure 3. HSMA can start as runtime software and evolve toward hardware-assisted semantic memory movement.

Component	Near-term role	Long-term hardware path	Why it matters
Semantic Memory Controller	Runtime policy engine for promotion/demotion	Hardware-assisted metadata and movement engine	Coordinates memory by semantic utility, not only address locality
Knowledge Cache	GPU/host-resident active concepts and subgraphs	On-package or near-memory semantic cache	Keeps active reasoning structure close to compute
KV Demotion Engine	GPU kernels for quantization, pruning, clustering	Dedicated compression/decompression path	Reduces HBM pressure and KV bandwidth
Graph Accelerator	Software graph traversal and ranking	Specialized traversal/ranking primitives	Accelerates concept navigation for retrieval
Reconstruction Engine	Replay, retrieve, or rehydrate context on demand	Batched context rebuilding engine	Avoids full-history reload while restoring needed detail
Memory Fabric	HBM, DRAM, CXL/pool, SSD tiers	Semantic-aware placement and movement	Supports long-lived agents with multiple fidelity tiers
KPU	Research concept only	Knowledge Processing Unit	Could accelerate concept merge/split, summary updates, graph maintenance

This aligns with broader hardware trends: high-bandwidth GPU fabrics such as NVLink target low-latency GPU-to-GPU communication, while CXL provides a cache-coherent interconnect for processors, accelerators, and memory expansion. HSMA would add a semantic placement policy above these physical fabrics [11, 12].

8. Why This Could Become a New Systems Layer for Agentic AI

Classical systems introduced abstractions that made larger applications possible: virtual memory made address spaces larger than physical memory; cache hierarchies hid latency; databases separated logical queries from physical storage.

HSMA analogy


HSMA proposes a comparable abstraction for agentic AI: semantic memory management separates the agent's logical understanding from the physical representation of tokens, KV, summaries, and graph state.

Classical systems layer	Problem solved	HSMA analogue
Virtual memory	Programs exceed physical memory	Agents exceed fixed context windows
Cache hierarchy	Hide latency and exploit locality	Keep active reasoning state close to compute
Database indexing	Retrieve by structure rather than scan	Retrieve by concept graph rather than transcript scan
Storage hierarchy	Move cold data to cheaper media	Demote cold context to summaries/graph/archive

If agentic workloads become dominant, memory management may become as strategically important as model inference throughput. The architecture that controls what the agent remembers, forgets, abstracts, and reconstructs could become a foundational platform layer.

9. Research Roadmap

Proposed Research Roadmap



Decision gates: quality retention, memory savings, reconstruction latency, hardware cost per recovered token.

Figure 4. Suggested staged research roadmap with decision gates.

Phase	Goal	Key experiments	Exit criteria
0. Baseline measurement	Quantify problem size	Measure KV footprint, bandwidth, recall failures, and long-agent latency	Clear bottleneck model and workload suite
1. Software prototype	Build semantic hierarchy	Extract summaries, facts, concept nodes, pointer maps	Improved long-horizon recall with controlled overhead
2. Reconstruction	Prove recoverability	Compare raw retrieval, partial replay, compressed KV rehydration	Target quality within tolerance at lower memory
3. Runtime policies	Optimize movement	Promotion/demotion policies using attention, reuse, semantic centrality	Stable policies with predictable latency
4. Hardware study	Evaluate acceleration	Simulate compression, graph traversal, reconstruction, interconnect movement	Cost/performance case for hardware investment

10. Success Metrics

Metric	Target direction	Why leadership should care
Peak KV memory	Reduce at fixed quality	Lowers HBM requirements and cost per agent
HBM bandwidth per generated token	Reduce	Improves inference throughput and utilization
Long-horizon recall accuracy	Maintain or improve	Prevents agents from losing earlier decisions
Selective reconstruction latency	Bound tightly	Avoids stalls when old context becomes relevant
Task completion on agent benchmarks	Improve or hold	Proves memory savings do not break reasoning
Joules per completed task	Reduce	Captures system-level value beyond tokens/sec

11. Technical Risks and Open Questions

Risk	Why it is hard	Mitigation / test
Semantic fidelity	Summaries and concepts may omit details that matter later	Use downward pointers, confidence scores, and reconstruction tests
Future utility prediction	Low-attention information may become important later	Combine attention, reuse history, task state, and graph centrality
Reconstruction quality	Regenerated context may differ from original neural state	Compare replay, raw text retrieval, compressed KV, and hybrid methods
Graph drift	Long-lived graphs may become inconsistent or overgrown	Use versioning, decay, merge/split operations, and audit trails
Hardware justification	Semantic operations may be too irregular for silicon	Start with profiling and target only stable, high-frequency kernels
Security and privacy	Persistent memory can retain sensitive data	Add governance, deletion, provenance, and policy-aware demotion

Critical validation requirement

The proposal should not proceed to hardware investment until a software prototype demonstrates measurable memory savings without unacceptable degradation in long-horizon reasoning, recall, and tool-use performance.

12. Investment Rationale

HSMA is not a near-term replacement for KV compression, RAG, or GraphRAG. It is a system architecture that can incorporate them and expose a new control plane for long-lived agents. The investment case is strongest if three conditions hold:

1. Agentic workloads become persistent enough that raw context retention is economically unattractive.
2. Semantic recall and structured memory become quality differentiators, not optional features.
3. A measurable subset of memory operations can be accelerated or co-designed with hardware and interconnects.

Under those conditions, HSMA offers a credible path toward platform differentiation: not just faster inference, but better memory architecture for agents that operate over long horizons.

13. Selected Prior Art and References

This proposal intentionally builds on existing work. The references below are included to help separate foundational prior art from the proposed HSMA contribution.

- [1] Packer et al., "MemGPT: Towards LLMs as Operating Systems," arXiv:2310.08560, 2023.
- [2] Rae et al., "Compressive Transformers for Long-Range Sequence Modelling," arXiv:1911.05507, 2019.
- [3] Lewis et al., "Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks," arXiv:2005.11401, 2020.
- [4] Borgeaud et al., "Improving language models by retrieving from trillions of tokens," arXiv:2112.04426, 2021.
- [5] Microsoft Research, "Project GraphRAG: LLM-Derived Knowledge Graphs," project page, 2024-present.
- [6] Zhang et al., "H2O: Heavy-Hitter Oracle for Efficient Generative Inference of Large Language Models," arXiv:2306.14048, 2023.
- [7] Xiao et al., "Efficient Streaming Language Models with Attention Sinks," arXiv:2309.17453, 2023.
- [8] Liu et al., "KIVI: A Tuning-Free Asymmetric 2bit Quantization for KV Cache," arXiv:2402.02750, 2024.
- [9] Li et al., "SnapKV: LLM Knows What You are Looking for Before Generation," arXiv:2404.14469, 2024.
- [10] Cai et al., "PyramidKV: Dynamic KV Cache Compression based on Pyramidal Information Funneling," arXiv:2406.02069, 2024.
- [11] Compute Express Link Consortium, "About CXL," official overview, accessed 2026.
- [12] NVIDIA, "NVLink and NVLink Switch," official overview, accessed 2026.

Bottom line

HSMA is a proposal for a new agentic AI memory layer: one that manages raw KV, compressed neural state, semantic summaries, concept nodes, and knowledge graphs as a single reconstructable hierarchy. The strategic opportunity is to make long-horizon agents scale with understanding, not merely with accumulated tokens.