


EXECUTION SETUP: CORE COMPONENTS

Slide 1

A four-layer workflow that turns benchmark tasks into observable agent runs.



SWE-bench	Provides realistic software-engineering tasks and expected fixes.
Deep Agents	Orchestrates planning, execution, patch generation, and review.
Dynamo	Packages requests and records runtime and tool observability.
SGLang	Executes inference and returns tool-call and response outputs.

Together, these components let us trace how a task moves from benchmark input to agent behavior.

CURRENT PROGRESS

Slide 2

The workflow is now observable end to end, and some runs complete through patch creation.

What Now Works

- **Patch creation observed:** Some tasks now complete through to actual patch creation.
- **Multi-phase flow visible:** Runs now follow planning, execution, patch generation, and review.
- **Execution is multi-step:** The execution phase now moves through multiple logical steps instead of a single-shot attempt.
- **Tool use confirmed:** Tool activity is visible by phase, so we can verify what the workflow actually did.

What We Are Improving

- **Completion quality:** Improve performance across more repos and task types.
- **Execution depth:** Reduce inspection-heavy runs that stop short of meaningful edits.
- **Model capability:** Test a larger model to improve reasoning and patch quality.

Bottom line: the workflow is now visible; model capability is the next major lever.

RUN SET AT A GLANCE

Slide 3

A compact snapshot of the current reported run set before we look at the per-run tables.

RUNS ANALYZED

27

Current reported run set

TOOL EVIDENCE

27 / 27

Every run shows recorded tool activity

PATCH GENERATION

25 / 27

Reached patch-generation phase

NON-EMPTY PATCHES

6

Runs that produced real diff output

BENCHMARK **SWE-bench**

MODEL **Qwen/Qwen2.5-Coder-7B-Instruct**

HARNESS **Deepagents**

TASK SUMMARY (1 OF 2)

BENCHMARK SWE-bench · MODEL Qwen/Qwen2.5-Coder-7B-Instruct · HARNESS Deepagents

RUN	REPO	TASK SUMMARY	EXPECTED ACTION	PATCH	BASE COMMIT
190251	NodeBB	Email Validation Status Not Handled Correctly in ACP and Confirmation Logic	fix validation logic	Yes	1e137b07
190500	qutebrowser	The `hide_qt_warning` function and its associated tests have been moved from `log.py` to `qtlog.py` to better organize Qt-specific logging functionality.	refactor code organization	Yes	ebfe9b7a
190534	NodeBB	Move .well-known assets to separate router file, add a basic webfinger implementation	modify routing/controller logic	Yes	da2441b9
190952	qutebrowser	Host blocking does not apply to subdomains when only the parent domain is listed	fix host-matching logic	Yes	0b8cc812
193512	NodeBB	Email Validation Status Not Handled Correctly in ACP and Confirmation Logic	fix validation logic	Yes	1e137b07
193615	qutebrowser	The `hide_qt_warning` function and its associated tests have been moved from `log.py` to `qtlog.py` to better organize Qt-specific logging functionality.	refactor code organization	Yes	ebfe9b7a
193649	NodeBB	Move .well-known assets to separate router file, add a basic webfinger implementation	modify routing/controller logic	Yes	da2441b9
193924	qutebrowser	Host blocking does not apply to subdomains when only the parent domain is listed	fix host-matching logic	Yes	0b8cc812
195404	NodeBB	Email Validation Status Not Handled Correctly in ACP and Confirmation Logic	fix validation logic	Yes	1e137b07
195450	qutebrowser	The `hide_qt_warning` function and its associated tests have been moved from `log.py` to `qtlog.py` to better organize Qt-specific logging functionality.	refactor code organization	Yes	ebfe9b7a
195524	NodeBB	Move .well-known assets to separate router file, add a basic webfinger implementation	modify routing/controller logic	Yes	da2441b9
195837	qutebrowser	Host blocking does not apply to subdomains when only the parent domain is listed	fix host-matching logic	Yes	0b8cc812
003009	NodeBB	Email Validation Status Not Handled Correctly in ACP and Confirmation Logic	fix validation logic	Yes	1e137b07
003113	qutebrowser	The `hide_qt_warning` function and its associated tests have been moved from `log.py` to `qtlog.py` to better organize Qt-specific logging functionality.	refactor code organization	Yes	ebfe9b7a

TASK SUMMARY (2 OF 2)

BENCHMARK SWE-bench · MODEL Qwen/Qwen2.5-Coder-7B-Instruct · HARNESS Deepagents

RUN	REPO	TASK SUMMARY	EXPECTED ACTION	PATCH	BASE COMMIT
003149	NodeBB	Move .well-known assets to separate router file, add a basic webfinger implementation	modify routing/controller logic	Yes	da2441b9
003318	qutebrowser	Host blocking does not apply to subdomains when only the parent domain is listed	fix host-matching logic	Yes	0b8cc812
003750	ansible	Collection Name Validation Accepts Python Keywords	fix validation logic	Yes	f533d465
010203	NodeBB	Email Validation Status Not Handled Correctly in ACP and Confirmation Logic	fix validation logic	Yes	1e137b07
010402	qutebrowser	The `hide_qt_warning` function and its associated tests have been moved from `log.py` to `qtlog.py` to better organize Qt-specific logging functionality.	refactor code organization	Yes	ebfe9b7a
010438	NodeBB	Move .well-known assets to separate router file, add a basic webfinger implementation	modify routing/controller logic	Yes	da2441b9
010608	qutebrowser	Host blocking does not apply to subdomains when only the parent domain is listed	fix host-matching logic	Yes	0b8cc812
011039	ansible	Collection Name Validation Accepts Python Keywords	fix validation logic	Yes	f533d465
125502	NodeBB	Email Validation Status Not Handled Correctly in ACP and Confirmation Logic	fix validation logic	Yes	1e137b07
125608	qutebrowser	The `hide_qt_warning` function and its associated tests have been moved from `log.py` to `qtlog.py` to better organize Qt-specific logging functionality.	refactor code organization	Yes	ebfe9b7a
125644	NodeBB	Move .well-known assets to separate router file, add a basic webfinger implementation	modify routing/controller logic	Yes	da2441b9
125814	qutebrowser	Host blocking does not apply to subdomains when only the parent domain is listed	fix host-matching logic	Yes	0b8cc812
130247	ansible	Collection Name Validation Accepts Python Keywords	fix validation logic	Yes	f533d465

RUN OVERVIEW (1 OF 2)

Slide 6

BENCHMARK SWE-bench · MODEL Qwen/Qwen2.5-Coder-7B-Instruct · HARNESS Deepagents

RUN	REPO	MODEL	STEPS	PLANNING	EXECUTION	PATCH GEN	REVIEW	OTHER	TOTAL	PATCH
190251	NodeBB	Qwen 2.5	3	0 · —	6 · read	2 · read	2 · read	4 · read	10	0 B
190500	qutebrowser	Qwen 2.5	3	0 · —	0 · —	2 · read	2 · read	4 · read	4	0 B
190534	NodeBB	Qwen 2.5	3	0 · —	5 · read	0 · —	0 · —	0 · —	5	0 B
190952	qutebrowser	Qwen 2.5	3	0 · —	0 · —	3 · read, todos	2 · read, todos	5 · read, todos	5	0 B
193512	NodeBB	Qwen 2.5	3	0 · —	0 · —	2 · exec, read	2 · exec, read	4 · exec, read	4	0 B
193615	qutebrowser	Qwen 2.5	3	0 · —	0 · —	2 · read	2 · read	4 · read	4	0 B
193649	NodeBB	Qwen 2.5	3	0 · —	7 · edit, exec, read, write	5 · exec	3 · exec	8 · exec	15	1.6 KB
193924	qutebrowser	Qwen 2.5	3	0 · —	0 · —	3 · read, todos	2 · read, todos	5 · read, todos	5	0 B
195404	NodeBB	Qwen 2.5	3	0 · —	0 · —	2 · exec, read	2 · exec, read	4 · exec, read	4	0 B
195450	qutebrowser	Qwen 2.5	3	0 · —	0 · —	2 · read	2 · read	4 · read	4	0 B
195524	NodeBB	Qwen 2.5	3	0 · —	5 · read	0 · —	3 · exec	3 · exec	8	0 B
195837	qutebrowser	Qwen 2.5	3	0 · —	0 · —	3 · read, todos	2 · read, todos	5 · read, todos	5	0 B
003009	NodeBB	Qwen 2.5	2	0 · —	5 · edit, exec, read	2 · write	1 · read	3 · read, write	8	412 B
003113	qutebrowser	Qwen 2.5	3	0 · —	0 · —	1 · read	2 · read	3 · read	3	0 B

Rows 1-14 of 27 · phase cells show tool-call count and dominant tools

RUN OVERVIEW (2 OF 2)

Slide 7

BENCHMARK SWE-bench · MODEL Qwen/Qwen2.5-Coder-7B-Instruct · HARNESS Deepagents

RUN	REPO	MODEL	STEPS	PLANNING	EXECUTION	PATCH GEN	REVIEW	OTHER	TOTAL	PATCH
003149	NodeBB	Qwen 2.5	2	0 · —	9 · edit, exec, read, write	1 · exec	1 · exec	2 · exec	11	1.9 KB
003318	qutebrowser	Qwen 2.5	3	0 · —	0 · —	1 · todos	1 · exec	2 · exec, todos	2	0 B
003750	ansible	Qwen 2.5	3	2 · todos	0 · —	1 · read	1 · read	4 · read, todos	4	0 B
010203	NodeBB	Qwen 2.5	3	0 · —	4 · read	5 · read	2 · read, todos	7 · read, todos	11	0 B
010402	qutebrowser	Qwen 2.5	3	0 · —	0 · —	1 · read	2 · read	3 · read	3	0 B
010438	NodeBB	Qwen 2.5	2	0 · —	9 · edit, exec, read, write	1 · exec	1 · exec	2 · exec	11	1.9 KB
010608	qutebrowser	Qwen 2.5	3	0 · —	0 · —	1 · todos	1 · exec	2 · exec, todos	2	0 B
011039	ansible	Qwen 2.5	3	2 · todos	0 · —	1 · read	1 · read	4 · read, todos	4	0 B
125502	NodeBB	Qwen 2.5	2	0 · —	5 · edit, exec, read	2 · write	1 · read	3 · read, write	8	412 B
125608	qutebrowser	Qwen 2.5	3	0 · —	0 · —	1 · read	2 · read	3 · read	3	0 B
125644	NodeBB	Qwen 2.5	2	0 · —	9 · edit, exec, read, write	1 · exec	1 · exec	2 · exec	11	1.9 KB
125814	qutebrowser	Qwen 2.5	3	0 · —	0 · —	1 · todos	1 · exec	2 · exec, todos	2	0 B
130247	ansible	Qwen 2.5	3	2 · todos	0 · —	1 · read	1 · read	4 · read, todos	4	0 B

Rows 15-27 of 27 · patch-producing runs are highlighted

WHAT THE RESULTS SHOW

Slide 8

What Is Already Working

- Multi-phase execution is now real and observable, not just conceptual.
- Every run shows tool-call evidence, and 16 runs already use write-oriented tools.
- A subset of runs now completes through patch creation, demonstrating true end-to-end workflow closure.

This means the orchestration stack can now produce inspectable evidence across the full task lifecycle.

What Needs The Next Push

- Patch-producing runs are currently concentrated in one repo family, so generality is still uneven.
- Many runs remain inspection-heavy, especially outside the strongest-performing task cases.
- The next lever is a larger model, measured against patch rate, repo coverage, and quality of final diffs.

Bottom line: the workflow is now observable end to end; the next step is improving task-solving capability.